

NAMAN YESHWANTH KUMAR

linkedin.com/in/namany | github.com/naman-ranka | nyeshwan@asu.edu | (623) 249-1265

SUMMARY

Full-stack engineer (MS Computer Engineering, ASU — GPA 3.81, May 2026) who builds reliable infrastructure for AI workloads end-to-end. Solo-shipped Trezzit (app.trezzit.com), a production Python + Django platform serving **600+ users** for 12+ months — async Celery job orchestration, Redis caching + rate limiting, multi-provider LLM routing (Gemini / OpenAI / Claude / Grok), per-endpoint p95 latency + per-provider token / cost telemetry. Built SiliconCrew, a LangGraph agent with async job engine (ThreadPoolExecutor, exponential backoff, on-disk run recovery) and full MCP server. Daily user of Cursor + Claude Code.

TECHNICAL SKILLS

Backend & APIs: Python | TypeScript | Django REST Framework | FastAPI | Node.js | PostgreSQL | Redis (caching + rate limiting) | Celery (async background jobs) | REST API Design | WebSocket | Cookie JWT + OAuth | Optimistic Locking

AI Infra & Job Orchestration: LangGraph | LangChain | MCP Protocol | Multi-provider LLM Routing (Gemini / OpenAI / Claude / Grok) | Async Job Engine (ThreadPoolExecutor, exponential backoff, on-disk recovery) | Per-endpoint p95 Latency + Per-provider Cost Telemetry | Cursor | Claude Code

Cloud, Frontend & OSS: AWS | Docker | DigitalOcean | Vercel | Supabase | Nginx | GitHub Actions CI/CD | Linux | Git | React 19 | Next.js | Tailwind | Intel reference architecture (PR #655 merged upstream)

PROJECT EXPERIENCE

Trezzit — Production Backend with Async Job Orchestration & Observability, Sole Engineer (app.trezzit.com) Feb 2025 – Present

- Solo-shipped Trezzit and run it on-call for **12+ months** — **60+ REST endpoints, 40 DB models, 995 tests** on PostgreSQL (Supabase), Redis (caching + rate limiting 100/hr anon, 10k/hr auth), Celery + Celery Beat + Flower for async background jobs; DigitalOcean VPS (prod) + AWS Lightsail (staging) + Vercel; GitHub Actions CI/CD with prod→staging DB sync and encrypted backups.
- Built a two-stage AI pipeline routing across Gemini / OpenAI / Claude / Grok — pass 1 + semantic validation gate + automatic model-upgrade fallback on failure — iterated parsing accuracy from **~60% → 95%+**; per-endpoint p95 latency tracking and per-provider token / cost telemetry across all 60+ endpoints — reliability + observability stack a small multi-cloud platform actually needs.
- Built an **MCP server** letting AI assistants take actions across the platform (create bills, query balances, upload receipts) over the same auth and data layer as the web app; cookie JWT with auto-refresh + failed-request queuing, Google OAuth, OTP, optimistic locking on edits.

SiliconCrew — LangGraph Agent with Async Job Engine & Multi-Provider LLM Routing Nov 2025 – Present

- Built a 21-tool LangGraph ReAct agent with **async job engine** (ThreadPoolExecutor, per-run isolation, exponential backoff 30s→10min, on-disk index for process-restart recovery); multi-provider LLM routing (Gemini / OpenAI / Anthropic) with provider-specific quirk handling; full **MCP server** (stdio / SSE / Streamable HTTP); achieved **37/92 on CVDP**; public at github.com/naman-ranka/siliconcrew.

HEAL.AI — AI Billing Assistant with RAG Pipeline (Devlabs Hackathon) Sep – Nov 2025

- Built HEAL.AI at Devlabs Hackathon — custom RAG pipeline (PDF chunking + 768-dim embeddings + cosine-similarity retrieval with confidence scoring) enabling natural-language Q&A over uploaded insurance policies and automated dispute-email generation via Gemini 2.5; FastAPI + Next.js + TypeScript, Dockerized.

Intel Automated Self-Checkout — Open Source Contribution Jan – Feb 2025

- Contributed sensor analytics microservice to **Intel's official self-checkout reference architecture** — configurable data pipeline (MQTT, Kafka, HTTP) for real-time LiDAR and weight sensor data with Grafana dashboard; **PR #655 merged upstream** after Intel engineering review.

WORK EXPERIENCE

Software Developer (Sole Engineer) Oct 2022 – Apr 2024
NCR GOLD Bengaluru, India

- Engineered a production Django platform (PostgreSQL + Celery + Nginx on DigitalOcean) processing **18,000+ orders over 3 years**; built 3 role-based portals (admin, artisan, customer) with multi-method auth — automated PDF invoice generation and WhatsApp Business API delivery on every state transition, eliminating **~20 min** of manual relay per order.
- Reverse-engineered third-party Jewely ERP (**zero public API**) through direct SQL schema analysis; built bidirectional bridge model syncing production weights and ERP status in real time; engineered inbound WhatsApp state-machine chatbot (JSON state per user, admin + customer flows) — full accountability loop across all three portals without disrupting legacy system.

EDUCATION

Master of Science — Computer Engineering (GPA 3.81/4.0) Expected May 2026
Arizona State University — Coursework: Artificial Intelligence · Foundations of Algorithms · Reinforcement Learning in Robotics · Tempe, AZ
Perception in Robotics · Algorithm/Hardware Co-Design

Bachelor of Engineering — Electrical and Electronics August 2022
BMS College of Engineering Bengaluru, India

AWARDS & HONORS

- Won **1st place at Devlabs Hackathon** with BillBusters (team), then **placed 2nd at the subsequent Devlabs Hackathon** solo-rebuilding it as HEAL.AI — consecutive hackathon placements iterating the same product under competition pressure.
- Won **'Best Use of AI'** at Strategy X DevHacks'25 (ASU) — built AdaptED AI, a Gemini-powered platform generating personalized flashcards, quizzes, and learning paths from uploaded documents in a 24-hour hackathon.
- Secured **2nd place at Hack SoDA 2024** (Amazon-sponsored, 24-hour) — built PassGen, a secure offline Chrome Extension for unique password generation.