NAMAN YESHWANTH KUMAR

linkedin.com/in/namany | github.com/naman-ranka | nyeshwan@asu.edu | (623) 249-1265

SUMMARY

MS Computer Engineering (ASU, GPA 3.81, May 2026) who builds production AI agents and iterates them to reliability. Sole-built Trezzit (trezzit.com), a live AI SaaS serving **600+ users** for 12+ months — iterated an LLM pipeline from **60% → 95%+** accuracy by diagnosing failure modes and building automated evaluation — and designed SiliconCrew, a 21-tool LangGraph agent with an autonomous eval-driven self-correction loop, scoring **37/92 on CVDP** benchmarks.

TECHNICAL SKILLS

Agents, Eval & RAG: LangGraph | ReAct Agents | MCP Protocol | Eval Pipelines | Autonomous Self-Correction | RAG Pipelines | Prompt Engineering | Multi-provider LLM Routing | Tool Use / Function Calling

Production Backend: Python | FastAPI | Django REST Framework | Async Python | PostgreSQL | Redis | Celery | Docker | GitHub Actions CI/CD | Linux

Frontend: React | Next.js | TypeScript | WebSocket | C++

EDUCATION

Master of Science — Computer Engineering (GPA 3.81/4.0) **Expected May 2026**
Arizona State University **Tempe, AZ**

Coursework: Artificial Intelligence · Reinforcement Learning in Robotics · Perception in Robotics · Foundations of Algorithms · Algorithm/Hardware Co-Design

Bachelor of Engineering — Electrical and Electronics **August 2022**
BMS College of Engineering **Bengaluru, India**

PROJECT EXPERIENCE

SiliconCrew — Production LLM Agent with Eval Framework & MCP Server **Nov 2025 – Present**

- Built a production LangGraph ReAct agent (21 tools) with an eval-driven self-correction loop — on task failure the agent reads diagnostic output, identifies root cause, patches its generation, and re-verifies without human input; achieved **37/92 first-pass** and **5/5 perfect score** on structured code generation benchmarks; a generalizable eval pattern applicable to any agent domain.
- Engineered multi-provider LLM routing across Gemini, OpenAI, and Anthropic with provider-specific quirk handling and quality-based fallback; built a full MCP server (stdio/SSE/HTTP transports, auto-discovery of tool functions) compatible with Claude Desktop, VS Code, and any MCP client; FastAPI + WebSocket backend with real-time tool-call streaming to a Next.js frontend.
- Built end-to-end observability: live streaming of tool-call execution traces, pipeline-stage artifact inspection at each agent step, and per-session token + cost tracking across all LLM providers — making agent behavior transparent and debuggable without re-running jobs.

Trezzit — Production AI SaaS (trezzit.com) **Feb 2025 – Present**

- Sole-built and shipped Trezzit (trezzit.com, **600+ users**, 12+ months live) — iterated receipt parsing from ~60% to **95%+** by diagnosing specific LLM failure modes (hallucinated item counts, currency ambiguity, multi-page layouts) and building a two-pass eval pipeline: extract structured data → semantically validate → escalate to a stronger model on failure; provider-abstracted registry across Gemini, OpenAI, Claude, and Grok.
- Production infrastructure: **995 tests** on every push via GitHub Actions CI with live PostgreSQL, Redis for caching and rate limiting, DigitalOcean VPS + Nginx + Gunicorn + systemd; **60+ REST endpoints, 40 Django models, 341 React components**; per-provider LLM cost and latency monitoring used to make ongoing model selection decisions.

HEAL.AI — Healthcare AI with Custom RAG Pipeline **Sep – Nov 2025**

- Won 1st place at Devlabs Hackathon; built a custom RAG pipeline without an external vector database — PDF chunking with 2-sentence overlap, Google text-embedding-004 (768-dim) embeddings, cosine similarity retrieval with confidence scoring, and automated FDCPA-compliant dispute email generation via Gemini; FastAPI backend, React + TypeScript frontend.

Intel Automated Self-Checkout — Open Source Contribution **Jan – Feb 2025**

- Contributed a real-time sensor data pipeline to **Intel's official open-source self-checkout reference architecture** — multi-transport ingestion layer (MQTT, Kafka, HTTP) for LiDAR and weight sensor streams, Grafana dashboard, and integration documentation; navigated an external production codebase, incorporated reviewer feedback, and had **PR #655 accepted by Intel engineers**.

WORK EXPERIENCE

Software Developer (Sole Engineer) **Oct 2022 – Apr 2024**
NCR GOLD **Bengaluru, India**

- Automated the entire order lifecycle for a wholesale business — production Django platform (PostgreSQL + Celery + Nginx) processing **18,000+ orders over 3 years**; reverse-engineered third-party ERP (zero public API) via SQL schema analysis for bidirectional sync, eliminating ~20 min of manual relay per order.

AWARDS & HONORS

- Won **1st place at Devlabs Hackathon** — HEAL.AI, an AI healthcare assistant with a custom RAG pipeline for insurance Q&A and automated medical bill dispute generation.
- Won **'Best Use of AI'** at Strategy X DevHacks'25 (ASU) — built AdaptED AI, a Gemini-powered platform generating personalized flashcards, quizzes, and learning paths from uploaded documents in a 24-hour hackathon.
- Secured **2nd place at Hack SoDA 2024** (Amazon-sponsored, 24-hour) — built PassGen, a secure offline Chrome Extension for password generation.